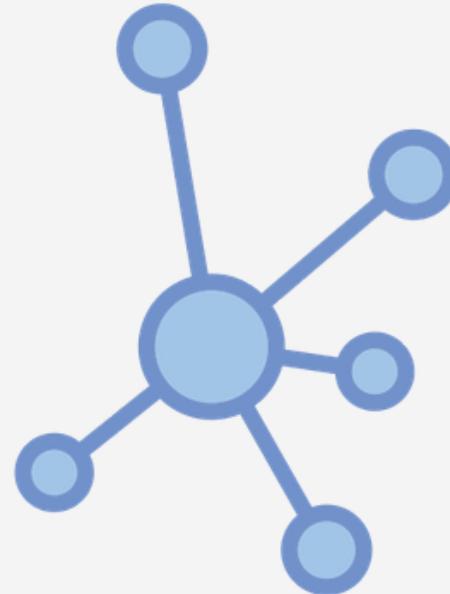




UNIVERSITÀ  
DI TRENTO

# Softwarized and Virtualized mobile networks

Network slice setup optimization



Dennis Cattoni  
Marco Lasagna  
Andrea Eugenio Cesaretti

# Network Slice Setup Optimization

***"GOAL:*** to enable RYU SDN controller to slice the network and then to dynamically re-allocate services in order to maintain desired QoS."

# Comnestsemu



- Comnetsemu repository:  
<https://git.comnets.net/public-repo/comnetsemu>

In the initial phase of development, we focused on understanding the **ComNetsEmu** environment by studying the codebase of the example scenarios presented during the lectures.

# Slices

The subsequent step focused on defining the **network slices**, the associated services, and the QoS policies. In particular, we identified two main slices:

1. Low-Latency Slice
2. High-Throughput Slice

# Low-Latency

**Goal:** The low latency slice is designed to minimize one-way data plane delay between a source host and a destination host within the SDN domain.

**Policy:**

- 30 ms threshold

**Features:**

- higher queue priority
- no service migration
- traffic simulated via *iperf3*

# High-Throughput

**Goal:** The high throughput slice is designed to maximize the sustained data rate achievable between a client host and a virtualized service endpoint within the SDN domain.

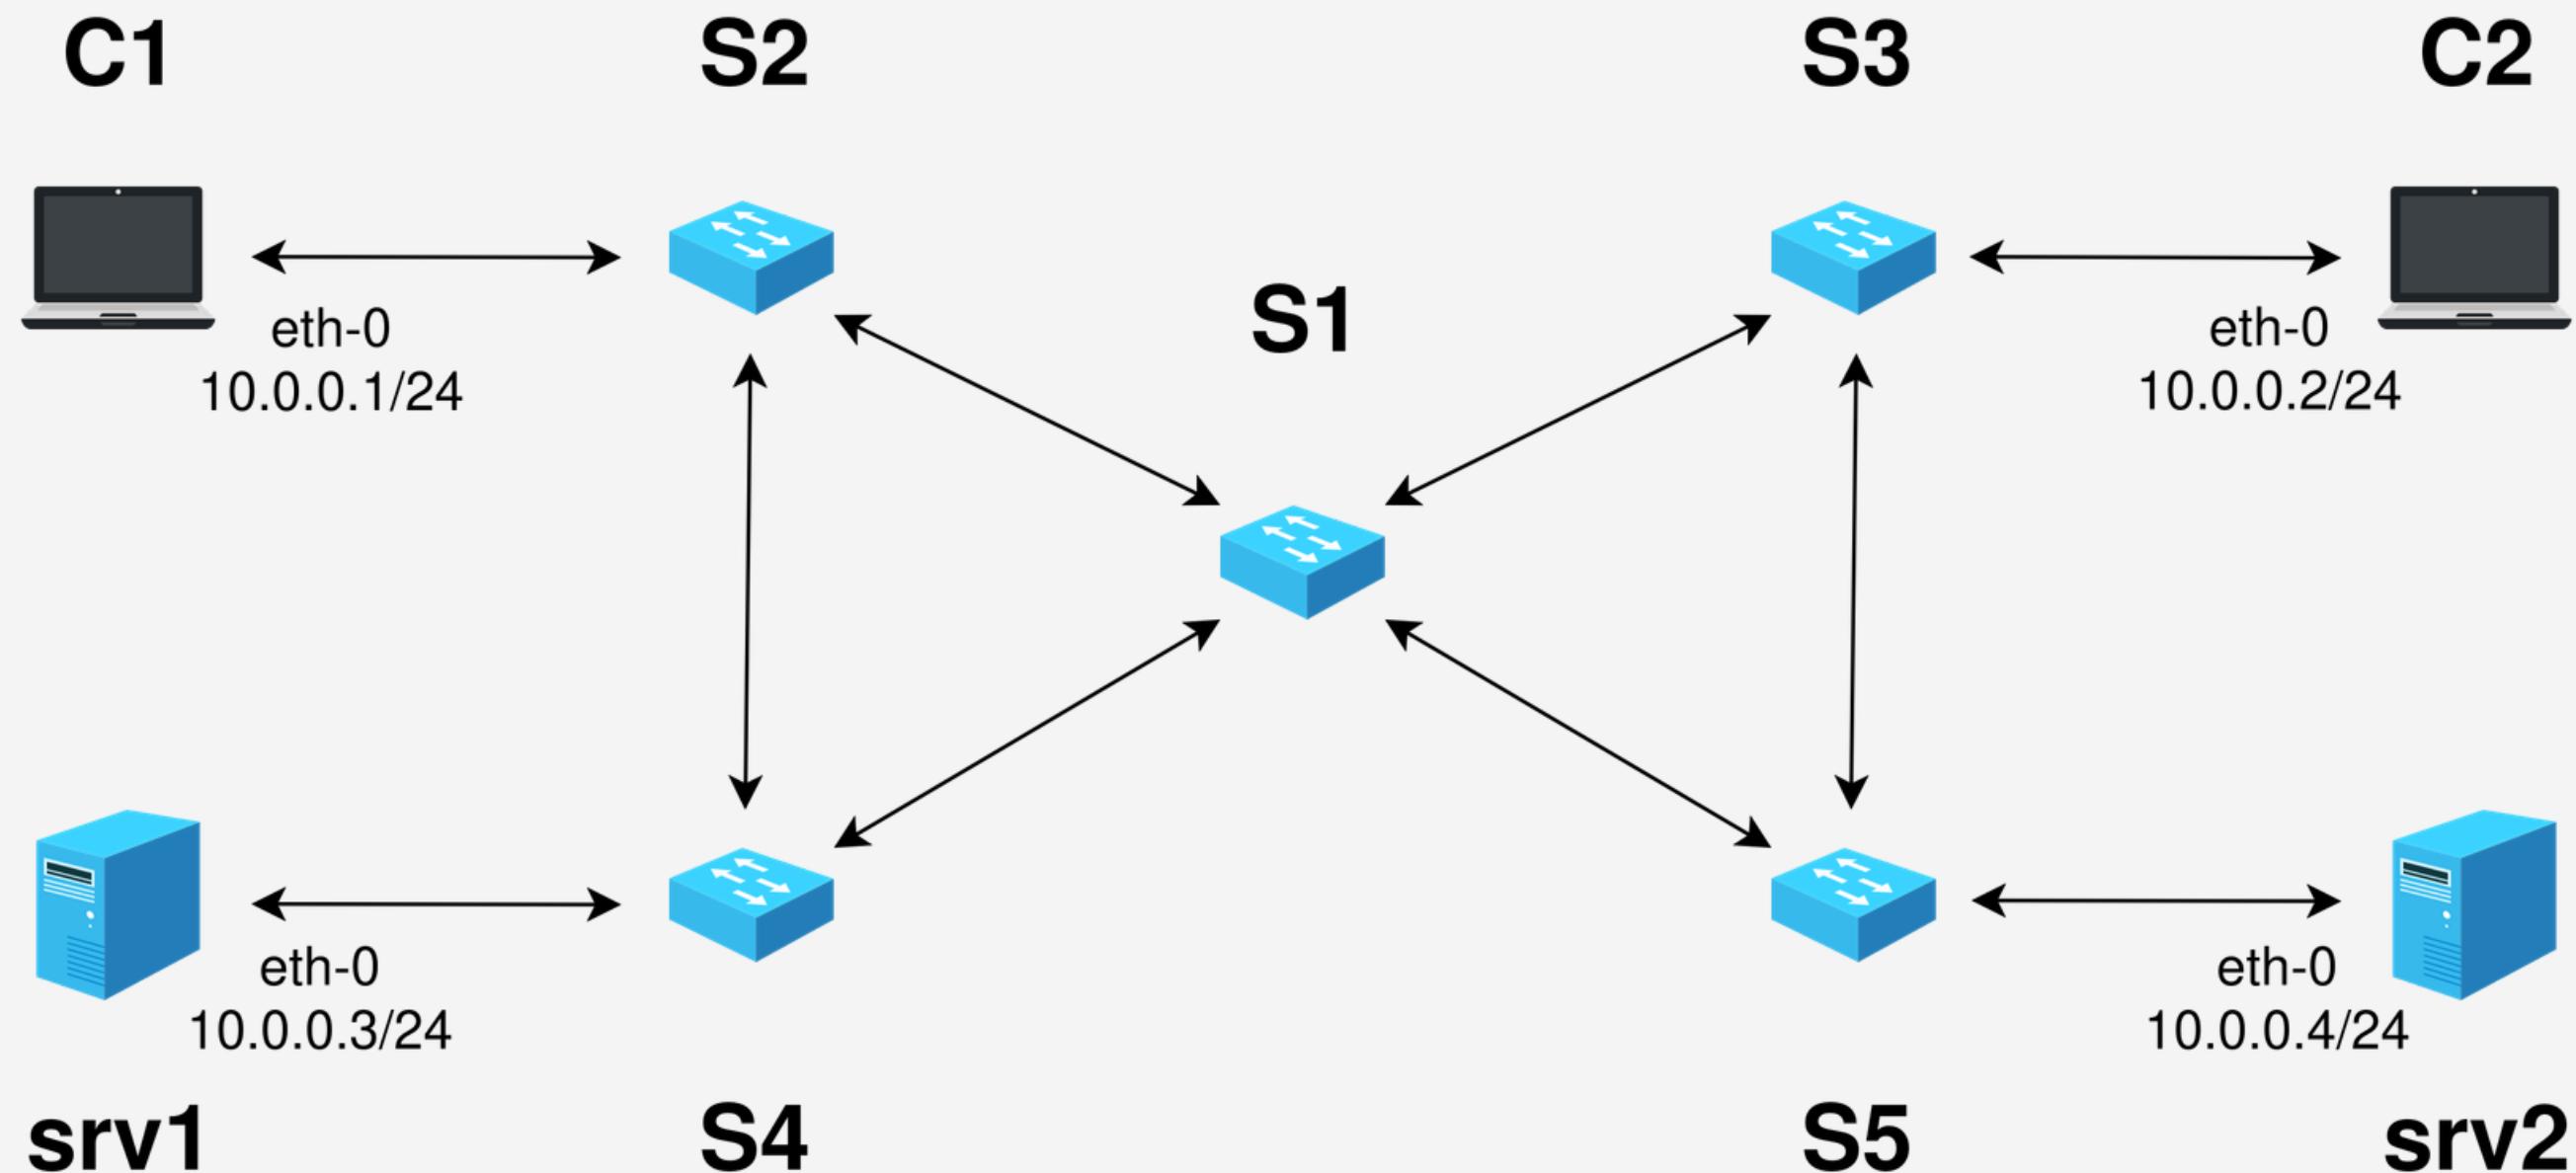
**Policy:**

- 6 Mbps threshold

**Features:**

- lower queue priority
- MEC service migration
- custom client and server

# Network topology



# OpenFlow



- Openflow Doc:  
<https://opennetworking.org/wp-content/uploads/2014/10/openflow-spec-v1.3.0.pdf>

The next step consisted of analyzing the characteristics of the **OpenFlow** southbound interface to identify which events, notifications, state changes, and operational information could be leveraged to design the overall optimization algorithm.

# Controller



- RYU controller:  
<https://ryu-sdn.org/>

Within the SDN architecture, **Ryu** operates as the SDN controller, implementing the control plane logic. The next step was to design and implement a custom Ryu controller from scratch in order to address the QoS optimization requirements of the defined network slices.

# Monitoring



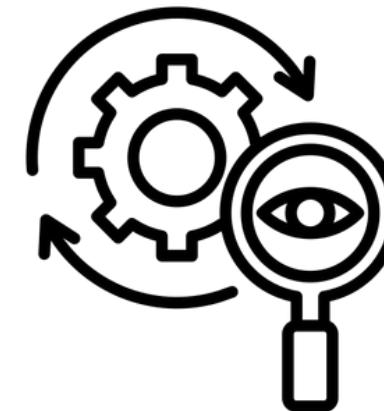
- QoS Monitoring with Dual-Tier Round-Robin

*(More details in the full spec)*

**Goal:** maintain an up-to-date network view for slice-aware QoS optimization, while balancing monitoring responsiveness and control-plane overhead.

# Monitoring

**Monitoring Model:** it relies on a periodic thread, which runs every  $\Delta T$  and collects per-port statistics via OpenFlow *PortStatsRequest* messages.

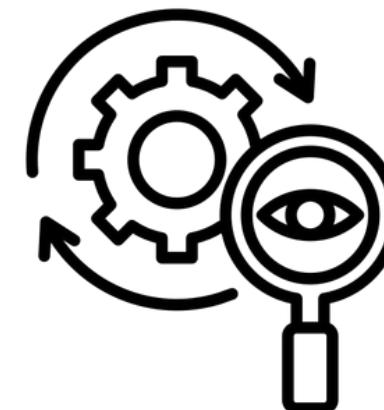


- **QoS Monitoring with Dual-Tier Round-Robin**  
*(More details in the full spec)*

# Monitoring

## Dual-Tier Strategy:

- Tier A – Active Slice Paths  
Poll all switches on active slice paths at every tick
- Tier B – Background Monitoring  
Poll K non-active switches per tick using RR



- **QoS Monitoring with Dual-Tier Round-Robin**  
*(More details in the full spec)*

# Optimization

# Low-Latency

It aims to continuously maintain a path that satisfies latency constraints by dynamically monitoring network conditions and triggering event-driven path recomputation upon link failures or QoS violations.

**Traffic:**

- simulated via *iperf3*

**Core feature:**

- Higher queue priority

# Low-Latency

1. Formally, the optimal path  $P^*$  is defined as:

$$P^* = \arg \min_{P:s \rightsquigarrow d} W_t(P)$$

2. The path cost is the sum of the costs associated with each link at time  $t$ :

$$W_t(P) = \sum_{i=0}^{n-1} w_t(v_i, v_{i+1})$$

3. The edge cost function is defined as:

$$w_t(u, v) = d(u, v) + \alpha \frac{\rho_t(u, v)}{1 - \rho_t(u, v) + \varepsilon} + p$$

# Low-Latency

4. The link utilization ratio between the estimated traffic rate on the link and its nominal capacity:

$$\rho_t(u, v) = \frac{\hat{x}_t(u, v)}{C(u, v)}$$

5. To reduce the impact of short-term fluctuations and measurement noise, the traffic rate estimate is obtained using an **EWMA**:

$$\hat{x}_t(u, v) = \beta \cdot x_t(u, v) + (1 - \beta) \cdot \hat{x}_{t-1}(u, v)$$

where:  $x_t(u, v) = \frac{\Delta B_t(u, v) \cdot 8}{\Delta t \cdot 10^6}$

# High-Throughput

It aims to maximize sustained data rate by dynamically monitoring path capacity, triggering event-driven path recomputation, and performing service migration when routing alone **cannot** satisfy throughput constraints.

## Traffic:

- Streaming on-demand

## Core features:

- Widest path
- MEC service migration

# High-Throughput

1. Formally, the optimal path  $P^*$  is defined as:

$$P^* = \arg \max_{P:s \leadsto d} B_t(P)$$

2. The bottleneck throughput is defined as the minimum residual capacity among all links composing the path:

$$B_t(P) = \min_{i=0, \dots, n-1} r_t(v_i, v_{i+1})$$

3. The residual capacity is computed as:

$$r_t(u, v) = C(u, v) - \hat{x}_t(u, v)$$

# High-Throughput

4. To reduce the impact of short-term fluctuations and measurement noise, the traffic rate estimate is obtained using an **EWMA**:

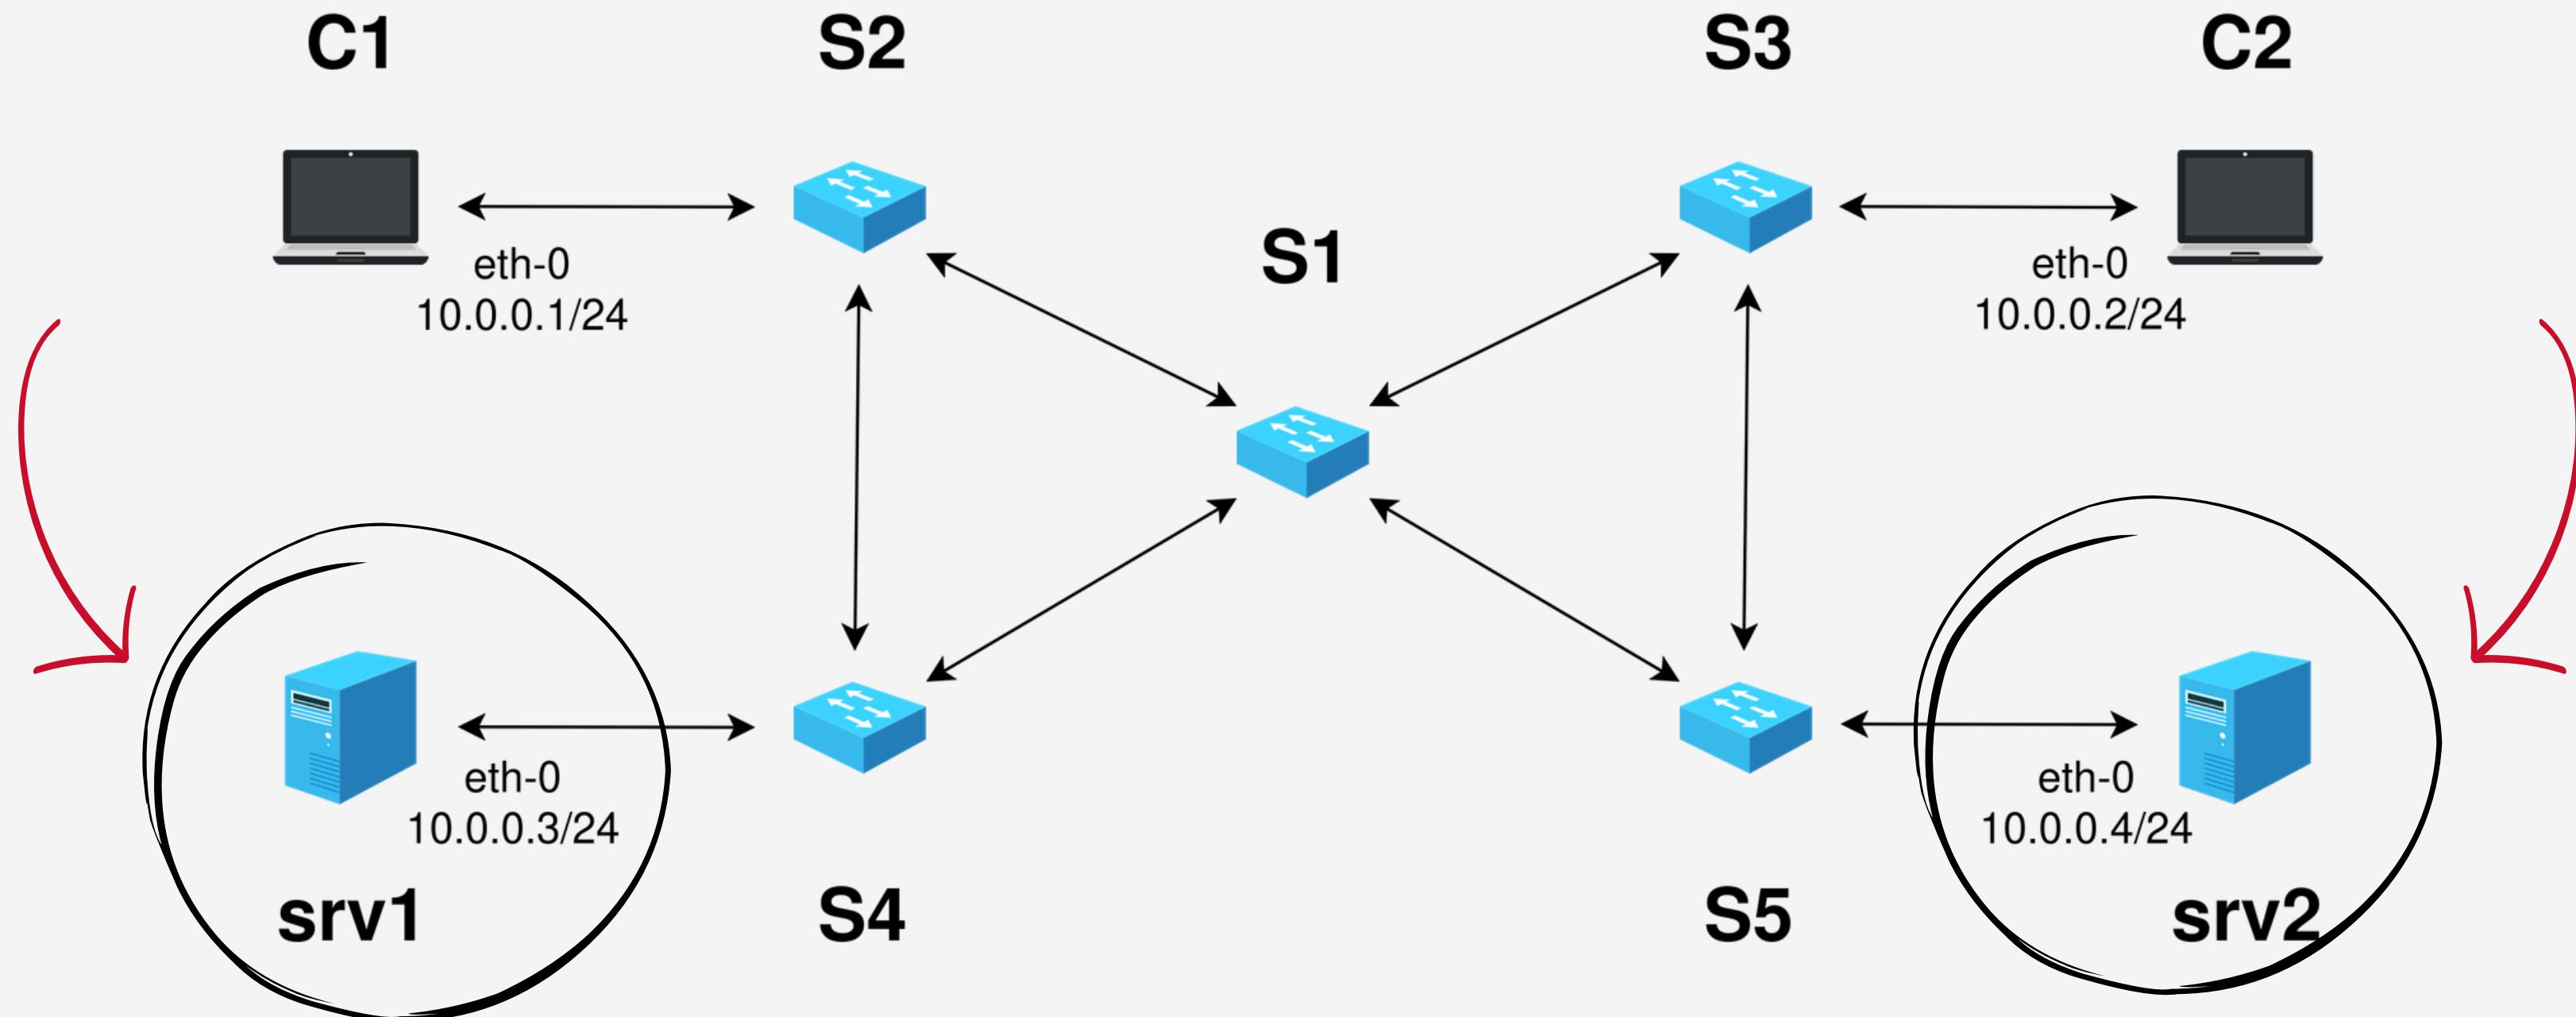
$$\hat{x}_t(u, v) = \beta \cdot x_t(u, v) + (1 - \beta) \cdot \hat{x}_{t-1}(u, v)$$

where:  $x_t(u, v) = \frac{\Delta B_t(u, v) \cdot 8}{\Delta t \cdot 10^6}$

# Service migration

**MEC** is leveraged as a complementary mechanism to SDN routing: when the throughput constraint cannot be satisfied through path recomputation, the *Orchestrator* instructs the *Controller* to migrate the service to an alternative backend on an edge node closer to the end user.

# Network topology

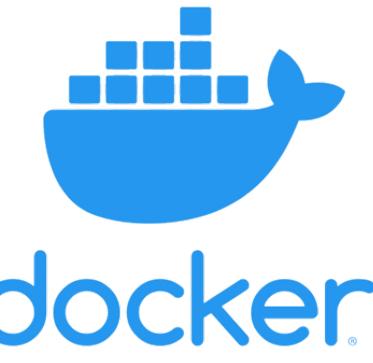




Almost  
done...

*Img source:*  
<https://deeprockgalactic.wiki.gg/wiki/APD-B317>

# NFV



- Docker:  
<https://www.docker.com/>

As a next step, **NFV** is used to virtualize the streaming on-demand service as a containerized network function, enabling controller-driven service migration that is invisible to the client, thanks to SDN-based VIP/VMAC and NAT.



Demo time!

# Limitations



## Known Limitations

*(More details in the full doc)*

### Environment Constraints:

the experimental setup is bound to ComNetsEmu  
on Ubuntu 20.04 LTS.

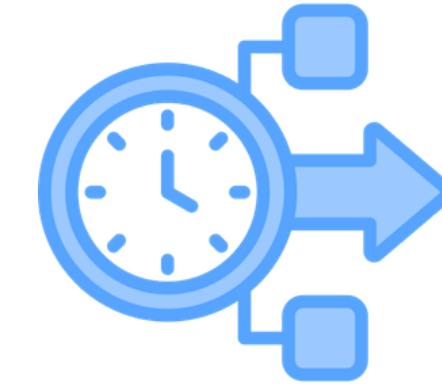
### Scalability Trade-offs:

balanced monitoring limits overhead but delays  
visibility on non-active links.

### Security:

it was not taken as a primary design goal.

# Future work



## Future work

(More details in the full doc)

Future work includes large-scale scalability evaluation, integration of modern transport protocols such as QUIC, and the incorporation of security-aware control policies and threat mitigation mechanisms.

*Thank you  
for your  
attention*



Marco Lasagna



Dennis Cattoni



Andrea Eugenio Cesaretti